

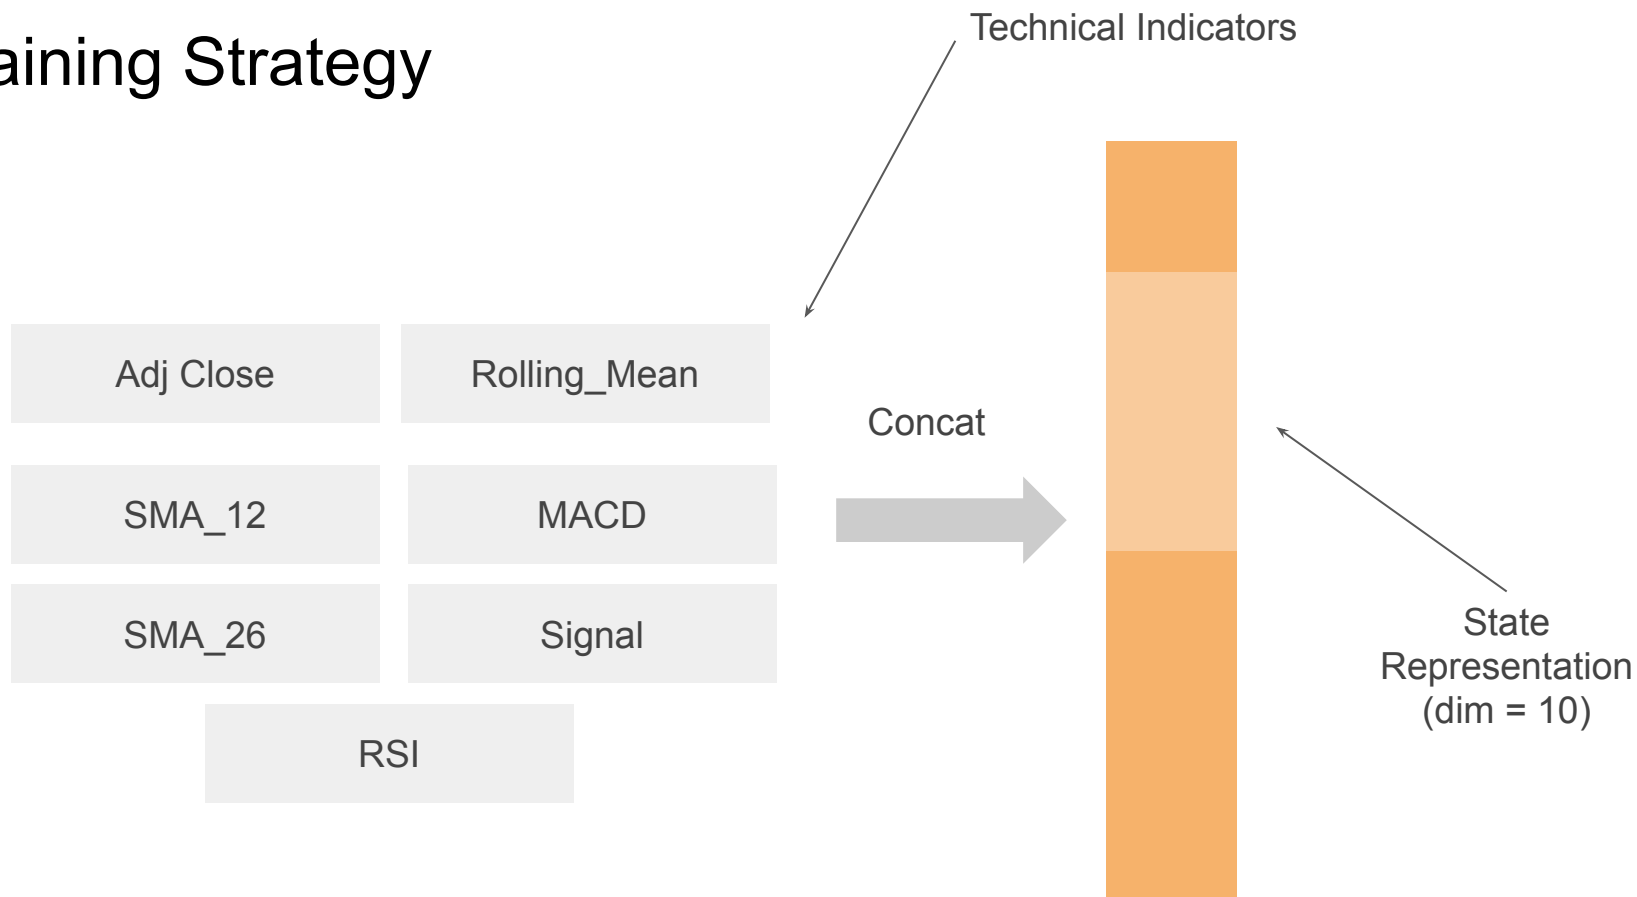
Fintech HW1 Repoert

r11944064 梁家綸

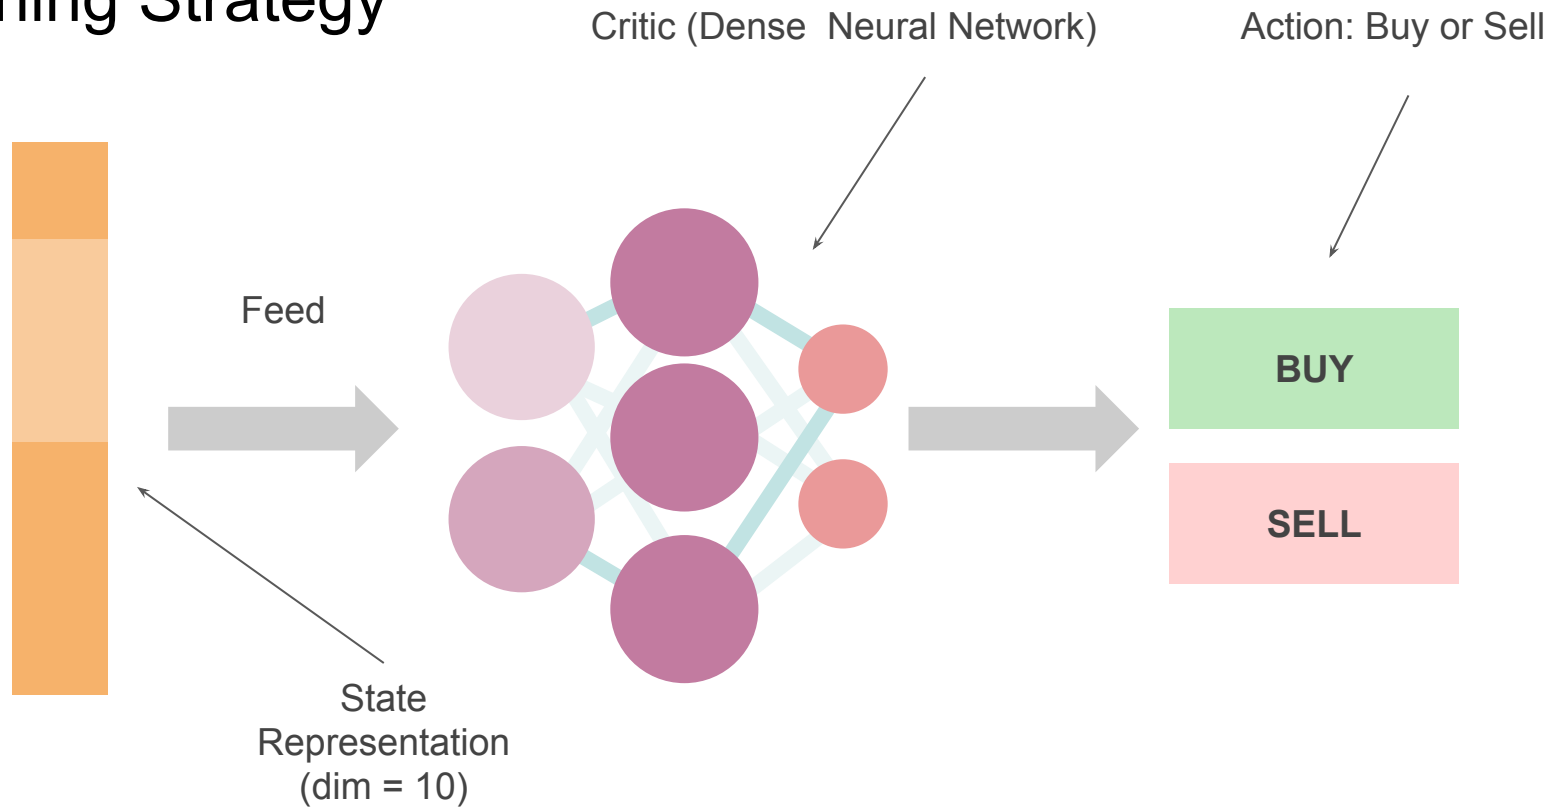
Summary

1. Github repository: https://github.com/b05505027/fintech_hw1
2. I trained a robo-advisor using **a DQN algorithm** with data from 0050.tw.
3. I tested the advisor on **the most recent 70 days of data** and trained it using **data from before those 70 days**.


Training Strategy




Training Strategy




Training Strategy



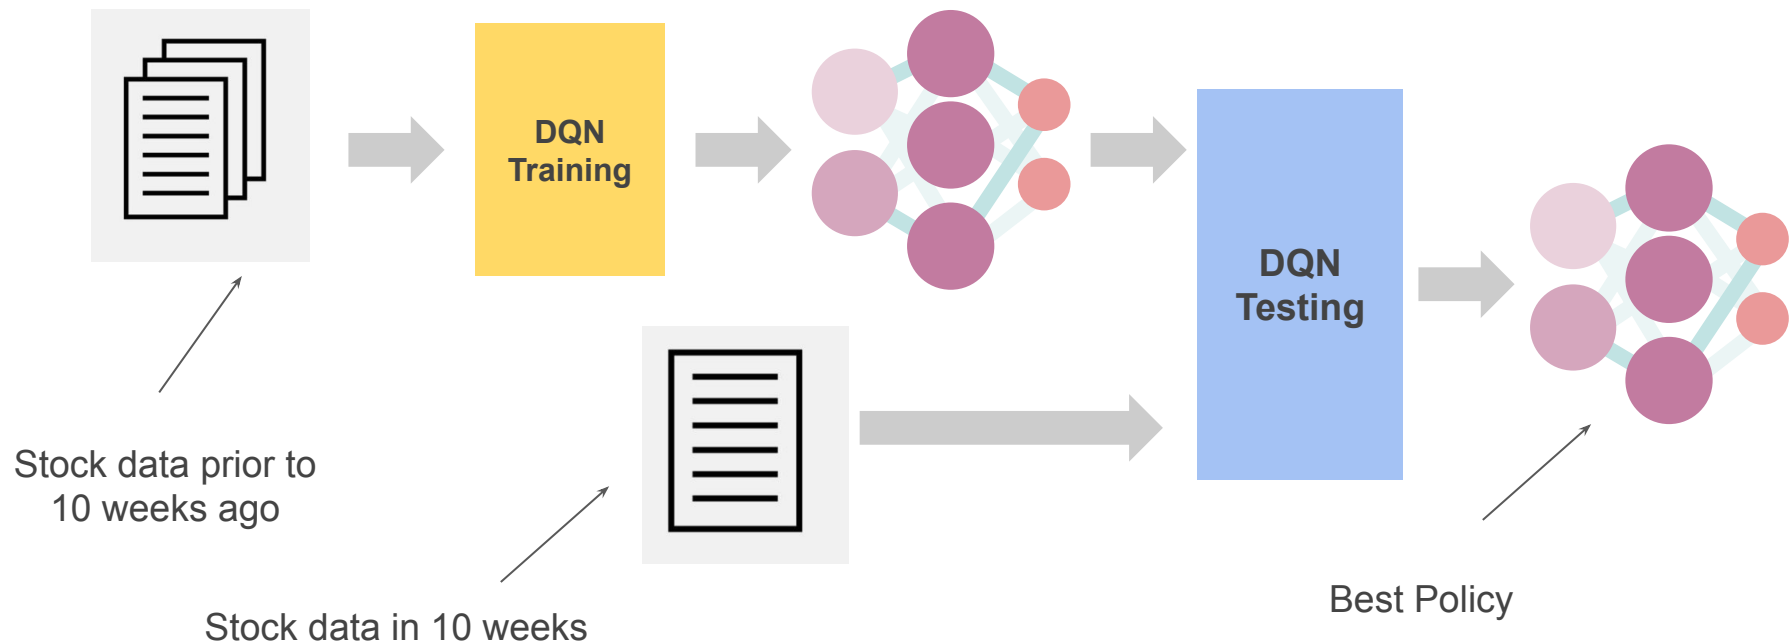
Training Strategy



DQN
Training Cycle



Training Strategy



Reward Function

Naturally, we aim to maximize the profit, and minimize the loss when holding stocks. So I choose the reward function:

$$r = 10 \cdot r_1 + 0.05 \cdot r_2$$

$r_1 = \sum_{t=1}^H \gamma^{t-1} \ln(G_t) \times \mathbb{1}[HoldingStock > 0]$	$G_t = \frac{StockPrice[t+1]}{StockPrice[t]}$
--	---

$r_2 = \sum_{t=1}^H \gamma^{t-1} HoldingStock \cdot (StockPrice_{t+1} - StockPrice_t)$
--

Algorithm

Algorithm 1 DQN Robo-Advisor


```
1: Initialize the critic network
2: Initialize the buffer_queue
3: exploration_rate  $\leftarrow$  1
4: for  $i = 0$  to episodes do
5:   randomly select a 70-days data from the historical data
6:   for  $j = 0$  to 69 do
7:     get the state
8:     if random_number  $<$  exploration_rate then
9:       use the random binary action
10:    else
11:      predict the action through the state and critic
12:    end if
13:    exploration_rate  $\leftarrow$  exploration_rate  $\times$  0.999
14:    based on the action, get the reward and the next state
15:    store transition  $(s, a, r, s')$  to the buffer_queue
16:    if  $j == 69$  then
17:      next_state  $\leftarrow$  None
18:    end if
19:  end for
20:  sample a batch of data from the buffer
21:  update the critic using MSELoss between  $r + Q(s', a')$  and  $Q(s, a)$ 
22: end for
23: return the best policy
```

Parameters


Table 1: DQN Robo-Advisor Parameters

Parameter	Value
Hidden Layer Size	64, 32 ; 400, 300
Input Size	7
Output Size	2
Normalization Method	Layer Normalization
Activation Function	ReLU
Optimizer	Adam
Learning Rate	1×10^{-3}
Training Episodes	100,000
Technical Indicators	Adj Close, MA12, SMA26, MACD, Signal, RSI, Rolling Mean
Batch Sizes	256, 128
Buffer Sizes	10000, 20000
Gamma	0.9, 0.95, 0.99


Training results: batch_size_128_buffer_size_5000_gamma_0.99




Training results: batch_size_128_buffer_size_5000_gamma_0.995




Training results: batch_size_128_buffer_size_5000_gamma_0.999




Training results: batch_size_256_buffer_size_10000_gamma_0.9



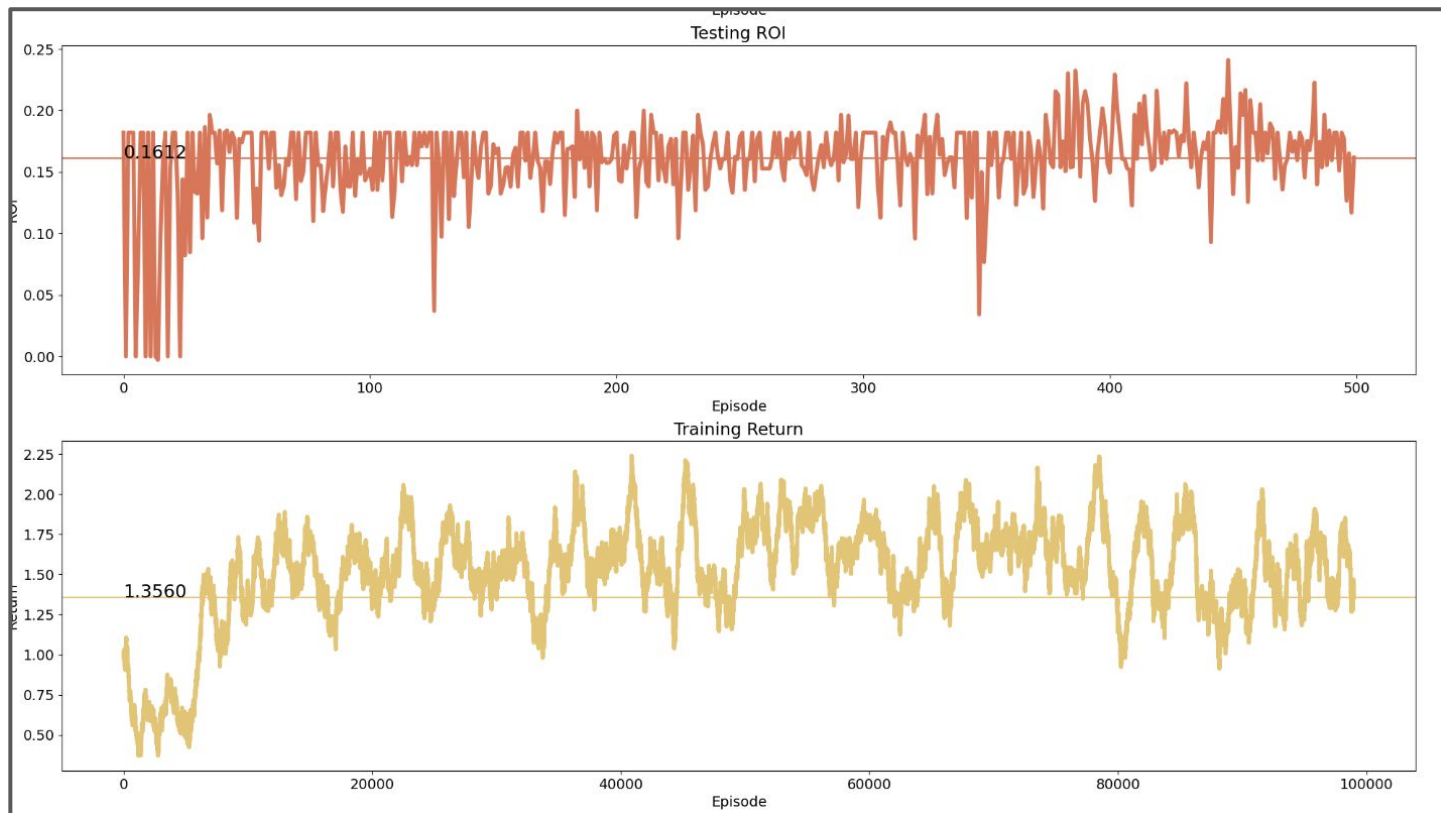
Training results: batch_size_256_buffer_size_10000_gamma_0.95




Training results: batch_size_256_buffer_size_10000_gamma_0.99



Training results: batch_size_256_buffer_size_20000_gamma_0.9



Training results: batch_size_256_buffer_size_20000_gamma_0.95



The Final Model

1. I choose the **batch_size_256_buffer_size_10000_gamma_0.95** as the final model.
2. For Simplicity, I didn't consider using the ensemble technique.
3. By executing the command `python rrEstimate.py 0050.TW-short.csv`, it can achieve **an approximate return rate of 110%**.

```
(fintech1) liangjialun@liangjialundeMacBook-Pro R11944064_t % python rrEstimate.py 0050.TW-short.csv  
rr=110.878849%
```